



(a) one σ for every x



(b) the network predicts $\sigma(x)$ too



(c) a second output head

