


(a) inference: each input is the model's previous output; the bold red token is where it goes wrong



(b) training with teacher forcing: gold inputs (amber); the losses $-\log p(\text{gold})$ above average to 0.168

